

CS100 Lecture 17

Rvalue References and Move

Contents

- Motivation: Copy is slow.
 - Rvalue references
- Move operations
 - Move constructor
 - Move assignment operator
 - The rule of five
- `std::move`
- NRVO, move and copy elision

Motivation: Copy is slow.

```
std::string a = some_value(), b = some_other_value();  
std::string s;  
s = a;  
s = a + b;
```

Consider the two assignments: `s = a` and `s = a + b`.

How is `s = a + b` evaluated?

Motivation: Copy is slow.

```
s = a + b;
```

1. Evaluate `a + b` and store the result in a temporary object, say `tmp`.
2. Perform the assignment `s = tmp`.
3. The temporary object `tmp` is no longer needed, hence destroyed by its destructor.

Can we make this faster?

Motivation: Copy is slow.

```
s = a + b;
```

1. Evaluate `a + b` and store the result in a temporary object, say `tmp`.
2. Perform the assignment `s = tmp`.
3. The temporary object `tmp` is no longer needed, hence destroyed by its destructor.

Can we make this faster?

- The assignment `s = tmp` is done by **copying** the contents of `tmp`?
- But `tmp` is about to "die"! Why can't we just *steal* the contents from it?

Motivation: Copy is slow.

Let's look at the other assignment:

```
s = a;
```

- **Copy** is necessary here, because `a` lives long. It is not destroyed immediately after this statement is executed.
- You cannot just "steal" the contents from `a`. The contents of `a` must be preserved.

Distinguish between the different kinds of assignments

```
s = a;
```

```
s = a + b;
```

What is the key difference between them?

- `s = a` is an assignment from an **lvalue**,
- while `s = a + b` is an assignment from an **rvalue**.

If we only have the copy assignment operator, there is no way of distinguishing them.

*** Define two different assignment operators, one accepting an lvalue and the other accepting an rvalue?**

Rvalue References

A kind of reference that is bound to **rvalues**:

```
int &r = 42;           // Error: Lvalue reference cannot be bound to rvalue.
int &&rr = 42;          // Correct: `rr` is an rvalue reference.
const int &cr = 42;     // Also correct:
                       // Lvalue reference-to-const can be bound to rvalue.
const int &&crr = 42;    // Correct, but useless:
                       // Rvalue reference-to-const is seldom used.

int i = 42;
int &&rr2 = i;           // Error: Rvalue reference cannot be bound to lvalue.
int &r2 = i * 42;        // Error: Lvalue reference cannot be bound to rvalue.
const int &cr2 = i * 42; // Correct
int &&rr3 = i * 42;      // Correct
```

- Lvalue references (to non-const) can only be bound to lvalues.
- Rvalue references can only be bound to rvalues.

Overload Resolution

Such overloading is allowed:

```
void fun(const std::string &);  
void fun(std::string &&);
```

- `fun(s1 + s2)` matches `fun(std::string &&)`, because `s1 + s2` is an rvalue.
- `fun(s)` matches `fun(const std::string &)`, because `s` is an lvalue.
- Note that if `fun(std::string &&)` does not exist, `fun(s1 + s2)` also matches `fun(const std::string &)`.

We will see how this kind of overloading benefit us soon.

Move Operations

Overview

The move constructor and the move assignment operator.

```
struct Widget {  
    Widget(Widget &&) noexcept;  
    Widget &operator=(Widget &&) noexcept;  
    // Compared to the copy constructor and the copy assignment operator:  
    Widget(const Widget &);  
    Widget &operator=(const Widget &);  
};
```

- Parameter type is **rvalue reference**, instead of lvalue reference-to- **const** .
- **noexcept** is (almost always) necessary! \Rightarrow We will talk about it in later lectures.

The Move Constructor

Take the `Dynarray` as an example.

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
public:
    Dynarray(const Dynarray &other) // copy constructor
        : m_storage(new int[other.m_length]), m_length(other.m_length) {
        for (std::size_t i = 0; i != m_length; ++i)
            m_storage[i] = other.m_storage[i];
    }
    Dynarray(Dynarray &&other) noexcept // move constructor
        : m_storage(other.m_storage), m_length(other.m_length) {
        other.m_storage = nullptr;
        other.m_length = 0;
    }
};
```

The Move Constructor

```
class Dynarray {  
    int *m_storage;  
    std::size_t m_length;  
public:  
    Dynarray(Dynarray &&other) noexcept // move constructor  
        : m_storage(other.m_storage), m_length(other.m_length) {  
  
    }  
};
```

1. *Steal* the resources of `other`, instead of making a copy.

The Move Constructor

```
class Dynarray {  
    int *m_storage;  
    std::size_t m_length;  
public:  
    Dynarray(Dynarray &&other) noexcept // move constructor  
        : m_storage(other.m_storage), m_length(other.m_length) {  
        other.m_storage = nullptr;  
        other.m_length = 0;  
    }  
};
```

1. *Steal* the resources of `other`, instead of making a copy.
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

* Take ownership of `other`'s resources!

The Move Assignment Operator

Take ownership of `other`'s resources!

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        m_storage = other.m_storage; m_length = other.m_length;  
  
        return *this;  
    }  
};
```

1. *Steal* the resources from `other`.

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        m_storage = other.m_storage; m_length = other.m_length;  
        other.m_storage = nullptr; other.m_length = 0;  
  
        return *this;  
    }  
};
```

1. *Steal* the resources from `other` .
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

Are we done?

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
  
        delete[] m_storage;  
        m_storage = other.m_storage; m_length = other.m_length;  
        other.m_storage = nullptr; other.m_length = 0;  
  
        return *this;  
    }  
};
```

0. Avoid memory leaks!

1. *Steal* the resources from `other`.
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

Are we done?

The Move Assignment Operator

```
class Dynarray {  
public:  
    Dynarray &operator=(Dynarray &&other) noexcept {  
        if (this != &other) {  
            delete[] m_storage;  
            m_storage = other.m_storage; m_length = other.m_length;  
            other.m_storage = nullptr; other.m_length = 0;  
        }  
        return *this;  
    }  
};
```

0. Avoid memory leaks!

1. *Steal* the resources from `other`.
2. Make sure `other` is in a valid state, so that it can be safely destroyed.

* Self-assignment safe!

Lvalues are Copied; Rvalues are Moved

Before we move on, let's define a function for demonstration.

Suppose we have a function that concatenates two `Dynarray` s:

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    for (std::size_t i = 0; i != a.size(); ++i)  
        result.at(i) = a.at(i);  
    for (std::size_t i = 0; i != b.size(); ++i)  
        result.at(a.size() + i) = b.at(i);  
    return result;  
}
```

Which assignment operator should be called?

```
a = concat(b, c);
```

Lvalues are Copied; Rvalues are Moved

Lvalues are copied; rvalues are moved ...

```
a = concat(b, c); // calls move assignment operator,  
                  // because `concat(b, c)` is an rvalue.  
a = b; // calls copy assignment operator
```

Lvalues are Copied; Rvalues are Moved

Lvalues are copied; rvalues are moved ...

```
a = concat(b, c); // calls move assignment operator,  
                  // because `concat(b, c)` generates an rvalue.  
a = b; // copy assignment operator
```

... but rvalues are copied if there is no move operation.

```
// If Dynarray has no move assignment operator, this is a copy assignment.  
a = concat(b, c)
```

Synthesized Move Operations

Like copy operations, we can use `=default` to require a synthesized move operation that has the default behaviors.

```
struct X {  
    X(X &&) = default;  
    X &operator=(X &&) = default;  
};
```

- The synthesized move operations perform member-wise initialization/assignment of the data members, in their declaration order, from an *xvalue*¹ argument.
 - i.e. `X(X&& o) : mem1(std::move(o.mem1)), mem2(std::move(o.mem2)), ... {}` .
(Later we will talk about `std::move` .)
- The synthesized move operations are `noexcept` if all the operations involved are `noexcept` .

Move operations can also be deleted by `=delete` , but be careful ...²

The Rule of Five: Idea

The updated *copy control members*:

- copy constructor
- copy assignment operator
- move constructor
- move assignment operator
- destructor

If one of them has a user-provided version, the copy control of the class is thought of to have special behaviors. (Recall "the rule of three".)

The Rule of Five: Rules

- The **move constructor** or the **move assignment operator** will not be generated if any of the rest four members have a user-declared version.
- The **copy constructor** or **copy assignment operator**, if not provided by the user, will be implicitly `delete`d if the class has a user-provided **move operation**.
- The generation of the **copy constructor** or **copy assignment operator** is **deprecated** (since C++11) when the class has a user-declared **copy operation** or a **destructor**.
 - This is why some of you see this error:

Implicitly-declared copy assignment operator is deprecated, because the class has a user-provided copy constructor.

The Rule of Five

The *copy control members* in modern C++:

- copy constructor
- copy assignment operator
- move constructor
- move assignment operator
- destructor

The Rule of Five: Define zero or five of them.

How to Invoke a Move Operation?

Suppose we give our `Dynarray` a label:

```
class Dynarray {  
    int *m_storage;  
    std::size_t m_length;  
    std::string m_label;  
};
```

The move assignment operator should invoke the **move assignment operator** on `m_label`. But how?

```
m_label = other.m_label; // calls copy assignment operator,  
                        // because `other.m_label` is an lvalue.
```

std::move

`std::move`

Defined in `<utility>`

`std::move(x)` performs an **lvalue to rvalue cast**:

```
int ival = 42;  
int &&rref = ival; // Error  
int &&rref2 = std::move(ival); // Correct
```

Calling `std::move(x)` tells the compiler that:

- `x` is an lvalue, but
- we want to treat `x` as an **rvalue**.

`std::move`

`std::move(x)` indicates that we want to treat `x` as an **rvalue**, which means that `x` will be *moved from*.

The call to `std::move` **promises** that we do not intend to use `x` again,

- except to assign to it or to destroy it.

A call to `std::move` is usually followed by a call to some function that moves the object, after which **we cannot make any assumptions about the value of the moved-from object**.

```
void foo(X &&x);           // moves `x`  
void foo(const X &x);      // copies `x`  
foo(std::move(x));        // matches `foo(X&&)` , so that `x` is moved.
```

"`std::move` does not *move* anything. It just makes a *promise*."

Use `std::move`

Suppose we give every `Dynarray` a special "label", which is a string.

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
    std::string m_label;
public:
    Dynarray(Dynarray &&other) noexcept
        : m_storage(other.m_storage), m_length(other.m_length),
          m_label(std::move(other.m_label)) { // !!
        other.m_storage = nullptr;
        other.m_length = 0;
    }
};
```

The standard library facilities ought to define efficient and correct move operations.

Use `std::move`

Suppose we give every `Dynarray` a special "label", which is a string.

```
class Dynarray {
    int *m_storage;
    std::size_t m_length;
    std::string m_label;
public:
    Dynarray &operator=(Dynarray &&other) noexcept {
        if (this != &other) {
            delete[] m_storage;
            m_storage = other.m_storage; m_length = other.m_length;
            m_label = std::move(other.m_label);
            other.m_storage = nullptr; other.m_length = 0;
        }
        return *this;
    }
};
```

The standard library facilities ought to define efficient and correct move operations.

Use `std::move`

Why do we need `std::move` ?

```
class Dynarray {  
public:  
    Dynarray(Dynarray &&other) noexcept  
        : m_storage(other.m_storage), m_length(other.m_length),  
          m_label(other.m_label) { // Isn't this correct?  
        other.m_storage = nullptr;  
        other.m_length = 0;  
    }  
};
```

`other` is an rvalue reference, so ... ?

An rvalue reference is an lvalue.

`other` is an rvalue reference, which is an lvalue.

- To move the object that the rvalue reference is bound to, we must call `std::move`.

```
class Dynarray {  
public:  
    Dynarray(Dynarray &&other) noexcept  
        : m_storage(other.m_storage), m_length(other.m_length),  
          m_label(other.m_label) { // `other.m_label` is copied, not moved.  
        other.m_storage = nullptr;  
        other.m_length = 0;  
    }  
};
```

An rvalue reference is an lvalue! Does that make sense?

Lvalues persist; Rvalues are ephemeral.

The lifetime of rvalues is often very short, compared to that of lvalues.

- Lvalues have persistent state, whereas rvalues are either **literals** or **temporary objects** created in the course of evaluating expressions.

An rvalue reference **extends** the lifetime of the rvalue that it is bound to.

```
std::string s1 = something(), s2 = some_other_thing();  
std::string &&rr = s1 + s2; // The state of the temporary object is "captured"  
                           // by the rvalue reference, without which the  
                           // temporary object will be destroyed.  
std::cout << rr << '\n'; // Now we can use `rr` just like a normal string.
```

Golden rule: **Anything that has a name is an lvalue.**

- The rvalue reference has a name, so it is an lvalue.

NRVO, Move and Copy Elision

Returning a Temporary (pure rvalue)

```
std::string foo(const std::string &a, const std::string &b) {  
    return a + b; // a temporary  
}  
std::string s = foo(a, b);
```

- First, a temporary is generated to store the result of `a + b`.
- How is this temporary returned?

Returning a Temporary (pure rvalue)

```
std::string foo(const std::string &a, const std::string &b) {  
    return a + b; // a temporary  
}  
std::string s = foo(a, b);
```

Since C++17, no copy or move is made here. The initialization of `s` is the same as

```
std::string s(a + b);
```

This is called **copy elision**.

Returning a Named Object

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    for (std::size_t i = 0; i != a.size(); ++i)  
        result.at(i) = a.at(i);  
    for (std::size_t i = 0; i != b.size(); ++i)  
        result.at(a.size() + i) = b.at(i);  
    return result;  
}  
a = concat(b, c);
```

- `result` is a local object of `concat`.
- Since C++11, `return result` performs a **move initialization** of a temporary object, say `tmp`.
- Then a **move assignment** to `a` is performed.

Named Return Value Optimization, NRVO

```
Dynarray concat(const Dynarray &a, const Dynarray &b) {  
    Dynarray result(a.size() + b.size());  
    // ...  
    return result;  
}  
Dynarray a = concat(b, c); // Initialization
```

NRVO transforms this code to

```
// Pseudo C++ code.  
void concat(Dynarray &result, const Dynarray &a, const Dynarray &b) {  
    // Pseudo C++ code. For demonstration only.  
    result.Dynarray::Dynarray(a.size() + b.size()); // construct in-place  
    // ...  
}  
Dynarray a@; // Uninitialized.  
concat(a@, b, c);
```

so that no copy or move is needed.

Named Return Value Optimization, NRVO

Note:

- NRVO was invented decades ago (even before C++98).
- NRVO is an **optimization**, but not mandatory.
- Even if NRVO is performed, the move constructor should still be available.
 - Because the compiler can choose not to perform NRVO.
 - The program should be syntactically correct ("well-formed"), no matter how the compiler treats it.

Summary

Rvalue references

- are bound to rvalues, and extends the lifetime of the rvalue.
- Functions accepting `x &&` and `const x &` can be overloaded.
- An rvalue reference is an lvalue.

Move operations

- take ownership of resources from the other object.
- After a move operation, the moved-from object should be in a valid state that can be safely assigned to or destroyed.
- `=default`
- The rule of five: Define zero or five of the special member functions.

Summary

`std::move`

- does not move anything. It only performs an lvalue-to-rvalue cast.
- `std::move(x)` makes a promise that `x` can be safely moved from.

In modern C++, unnecessary copies are greatly avoided by:

- copy-elision, which avoids the move or copy of temporary objects, and
- move, with the `return` ed lvalue treated as an rvalue, and
- NRVO, which constructs in-place the object to be initialized.

Notes

¹ An xvalue is an rvalue that has an identity. A typical xvalue is one obtained from `std::move` applied to an lvalue.

² We seldom delete move operations. In most cases, we want rvalues to be copied if move is not possible. An explicitly deleted move operation will make rvalues not copyable, because deleted functions still participate in overload resolution.

An implicitly deleted move operation will be ignored in overload resolution. This is addressed by the defect report [CWG 1402](#) and is applied retroactively to C++11. Note that this change of behavior did not come into effect when the book *C++ Primer, 5e* was published.